

The ClimaHealth Layer

Enterprise Product Specification, Process Flowcharts & Computational Architecture

1. Executive Summary & Market Mechanics


The ClimaHealth Layer is an enterprise-grade B2B predictive analytics middleware platform designed specifically for the United States healthcare ecosystem. It acts as an interoperable data-processing pipeline that connects real-time, hyper-local spatiotemporal environmental feeds with longitudinal electronic health record (EHR) registries. The platform's core commercial objective is to predict and mitigate acute clinical escalations across an entire population panel of unhospitalized, "at-home" chronic patients 24 to 72 hours before they experience a medical emergency.


In the US healthcare market, the financial burden of emergency care is heavily concentrated onto risk-bearing entities under **Value-Based Care (VBC)** models, **Accountable Care Organizations (ACOs)**, and **Medicare Advantage** infrastructure. When a stable chronic patient experiences an acute health crisis driven by an external climate catalyst (such as a severe ozone spike, wildfire smoke plume, or extreme heatwave), they flood emergency departments. These avoidable visits impose steep financial penalties and uncompensated care costs on health networks.

By moving beyond legacy, reactive 30-day post-discharge readmission trackers and continuously evaluating the ecological vulnerabilities of the broad, unhospitalized population, The ClimaHealth Layer protects enterprise health margins while establishing an automated, predictive population safety net.

2. System Process Flowchart

The architectural journey of data from initial ingestion through computational processing, clinical triage, and automated endpoint notification executes across a strictly synchronized, event-driven pipeline:





3. Microenvironment Exposure Modeling & The Core Mathematical Formulas

A foundational failure point of environmental tracking platforms is the reliance on raw macro-meteorological forecasts. If a system fires warnings based purely on regional coordinates, it triggers severe false positives for individuals sitting inside modern, climate-controlled, air-purified office structures. To decouple outdoor weather anomalies from true individual physiological stress, the platform executes a multi-tiered **Microenvironment Exposure Function**.

3.1 Static Infrastructure Filtering

The system assigns an unchanging, location-specific **Shielding Coefficient (SC)** to each patient profile by processing census-tract Social Determinants of Health (SDOH) datasets, neighborhood tree canopy coverage, building material metadata, and historical regional HVAC implementation density. A score approaching 1.0 represents highly insulated, modern structural resilience.

3.2 Dynamic Hardware Proxies

To determine a patient's active localization without compromising digital privacy, the consumer application monitors privacy-compliant hardware proxies on the smartphone edge. It verifies if the device is connected to the cryptographic hash of the home Wi-Fi signature, tracks pedometer movement, and queries the internal barometric sensor. If the internal pressure reading is flat and stable (indicating a pressurized, climate-controlled indoor HVAC setup) and the satellite GPS signal strength exhibits heavy roof attenuation, the backend validates that the patient is indoors.

3.3 The Core Exposure Attenuation Formula

When the edge indicators confirm a patient is indoors, the data ingestion pipeline intercepts the continuous climate vector before it reaches the deep learning core. It dynamically applies a mathematical multiplier to scale down the outdoor hazard weight, ensuring the predictive layers evaluate actual, high-fidelity physiological strain rather than generic outdoor numbers:

$$\text{Effective Exposure} = \text{Outdoor Hazard Concentration} \times (1 - \text{Shielding Coefficient})$$

3.4 Multi-Task Loss Optimization

Once the attenuated exposure parameters are calculated, they are passed into a unified **Temporal Fusion Transformer (TFT)** network. The model utilizes shared hidden temporal layers to find hidden biological cross-strain variables before splitting into distinct predictive heads. The entire backend network optimizes an integrated multi-task loss function during training:

$$L_{total} = \alpha L_{respiratory} + \beta L_{cardiovascular} + \gamma L_{metabolic}$$

Plain-English Breakdown of the Multi-Task Loss Equation

In artificial intelligence, "**Loss**" (represented by **L**) simply translates to "**Mistakes**." When an AI model is learning, its loss function calculates mathematically how incorrect its predictions are compared to real-world outcomes. The computational goal during training is to drive the total loss down as close to zero as possible.

Because **The ClimaHealth Layer** is a Multi-Task AI network, it must calculate and grade its own performance across three entirely independent chronic disease vectors at the same time. The individual parameters of the equation function as follows:

- **L_{total} (Total Loss):** The overall "Report Card" penalty score. It represents the combined sum of all mistakes the model has made across all tracking categories. The system constantly modifies its internal math layers to shrink this specific total weight.
- **L_{respiratory} (Respiratory Loss):** The specific count of *Lung Mistakes* made by the network (i.e., tracking how incorrect the model's predictions were regarding asthma or COPD flare-ups).
- **L_{cardiovascular} (Cardiovascular Loss):** The specific count of *Heart Mistakes* made by the network (i.e., tracking how incorrect the model's predictions were regarding cardiac strain or stroke risks).
- **L_{metabolic} (Metabolic Loss):** The specific count of *Blood Sugar Mistakes* made by the network (i.e., tracking how incorrect the model's predictions were regarding diabetic dehydration or glycemic crises).
- **α, β, γ (Alpha, Beta, Gamma Priority Knobs):** These symbols function as specialized mathematical "multiplier dials" or volume knobs. Since certain medical errors are vastly more dangerous than others (e.g., missing an acute heart attack is immediately fatal compared to a mild allergy warning), engineers adjust these multipliers to prioritize the AI's training focus. Turning up **β (Beta)** tells the algorithm that a heart mistake carries 5 times the penalty score of a lung mistake, forcing the neural network to focus its mathematical optimizations on minimizing cardiovascular failures.

4. Core Technical Feature Specifications

- **Multi-Task Prediction Heads:** Individual layers output specific risk probabilities: an **Airway Hyperresponsiveness Head** (Asthma/COPD flaring), an **Autonomic Cardiovascular Strain Head** (Stroke/Heart failure fluid tracking), and a **Thermoregulatory Glycemic Head** (Dehydration-induced metabolic failures).
- **Uber H3 Spatial Hexagonal Ingestion:** Converts floating-point latitudes and longitudes into unique 64-bit integer H3 hexagonal strings at **Resolution 7 or 8**. This converts costly geometric polygon intersection lookups into instantaneous key-value database matching, avoiding systemic database processing blocks.
- **72-Hour Provider Triage Dashboard:** Implements a visual registry for medical care coordinators. Patients are sorted dynamically by their **Climate Volatility Delta** (the explicit increase in crisis risk driven by oncoming weather anomalies). To satisfy clinical validation and safety needs, every flag embeds an Explainable AI (XAI) breakdown powered by SHAP core libraries.
- **48-Hour "Approve & Release" Workflow:** Generates tailored, non-prescriptive, white-labeled automated SMS/IVR campaigns. Care managers review pre-compiled text sheets and clear the whole population tier with a single click (`[Approve and Schedule Campaign]`), maintaining clinical oversight.

5. Defensive Engineering & Risk Mitigation Matrix

Identified Failure Point	The Operational / Clinical Threat	Defensive Product Engineering Resolution
FDA Medical Device (SaMD) Tripwire	Algorithms outputting real-time alarms for life-threatening events are regulated as high-risk medical devices, forcing years of clinical validation trials.	72-Hour Workflow Scheduling Window: Position the platform as a mid-range population planning tool. Moving predictions out to 3 days avoids real-time intervention categories, while transparent SHAP maps satisfy non-device Clinical Decision Support (CDS) parameters.
Extreme Clinician Alert Fatigue	Regional environmental shifts affect large geographic zones, triggering thousands of simultaneous patient flags that cause doctors to mute the system.	Token Bucket Triage Constrainer: Enforce constrained optimization algorithms that analyze a clinic's operational capacity and limit total daily warnings to the top 2% to 5% of highest-risk actionable deltas.
EHR "Swivel-Chair" Workflow Rejection	Forcing medical practitioners to leave their primary interface to log into an external website or separate software ensures zero clinical adoption.	SMART on FHIR Embedding: Wrap the entire user interface layer in the federal SMART on FHIR standard, embedding your platform natively as a secure application window right inside the Epic or Cerner charting screens.
Provider Legal Malpractice Liability	Flagging a patient as high-risk creates a permanent data record. If a busy doctor overlooks a notification, they assume profound liability if an emergency occurs.	FHIR Progress Note Ledger Encapsulation: Tie all alert boundaries strictly to peer-reviewed guidelines. When approved, the system fires a FHIR API write-back payload that documents the preventative outreach directly inside the native EHR chart as a legal shield.
Multi-Hardware Device Sparsity	Relying on streaming wearable telemetry means the predictive models will fail or error out if a patient forgets to charge or wear their hardware.	Hierarchical Graceful Pipeline Degradation: If streaming telemetry shows null values (\$NaN\$), the data ingestion workers drop wearable metrics and shift processing to an ensemble fallback layer using stable EHR records and climate feeds.
Patient Consumer "Anxiety Churn"	Direct-to-consumer clinical warnings cause acute psychological stress, driving immediate app uninstalls or uncoordinated panics to clinic phone banks.	Ambient Health-Weather Framing: Design the patient portal as a premium lifestyle daily weather utility. Present risk data as a positive "ClimaHealth Index" out of 100 with actionable daily behavioral nudges.

6. Interoperability & Regulatory Standards: SMART on FHIR

To achieve seamless, universal deployment across the United States healthcare ecosystem without requiring unique, custom codebases for individual hospital networks, the system architecture conformed strictly to the federally mandated **SMART on FHIR** framework.

FHIR: The Unified Data Protocol

Fast Healthcare Interoperability Resources (FHIR) replaces non-standardized, siloed legacy health formats with a universal JSON-based data protocol. Medical information is structured into highly predictable, standardized building blocks called *Resources*. The ClimaHealth Layer queries the EHR for clean, repeatable arrays such as the Patient identity, the Observation block (for clinical lab histories and historical vitals tracking), and the MedicationRequest resource (for prescription profiles).

SMART: The Identity & Integration Shield

SMART adds a robust identity verification layer over the FHIR data pipeline using the open OAuth 2.0 framework. It acts as the secure connection layer that allows the hospital system to log the provider into your platform using their native credentials, securely confirm exactly which patient record is active on their terminal, and project your visual dashboard as a clean, integrated frame directly inside the doctor's workflow window.

7. B2B Population Consent Architecture

Compliance with HIPAA and the Telephone Consumer Protection Act (TCPA) dictates that a patient must provide an explicit digital signature before a company can gather live geospatial telemetry or fire automated smartphone text messages. The platform implements a **Dual-Track Compliance Engine** to ensure full population coverage even when patients opt out:

- **Track A — Fully Consented Population (Automated Pipeline):** Activated when a patient signs the digital privacy disclosure on the consumer Health-Weather app. The system tracking engine maps their background H3 hexagonal coordinates, ingests active smart device vitals, and unlocks complete 48-hour automated SMS preventative outreach campaigns on autopilot.
- **Track B — Non-Consented Population (Internal Triage Only):** Activated when a patient declines application installation or blocks device location tracking. **Under the HIPAA Privacy Rule's "Healthcare Operations" provision, a health system is legally permitted to process a patient's medical records and home address for population care coordination without explicit opt-in.** Non-consented patients still securely map to the internal 72-Hour Provider Triage Dashboard based on their billing ZIP code. Care coordinators see the flag and execute a traditional, compliant, manual telephone call to deliver preventative care reminders, maintaining 100% platform utility.

8. Complete Production Technical Stack

Infrastructure Layer	Selected Enterprise Framework / Technology	Strategic Architectural Rationale
Ingestion & Streaming	Apache Kafka / AWS Kinesis Data Streams	Manages high-throughput, asynchronous, multi-rate streaming data buses without locking backend processing threads.
Data Interoperability	HL7 FHIR v4 / v5 (Google Cloud Healthcare API framework)	Deploys event-driven FHIR Subscriptions to instantly ingest clinical baseline shifts pushed directly from hospital servers.
Geospatial Indexing	Uber H3 (Hexagonal Hierarchical Spatial Indexing Core)	Replaces expensive trigonometric polygon calculations with sub-millisecond, integer-based string lookups.
Machine Learning Core	PyTorch Forecasting / Temporal Fusion Transformers (TFT)	Binds static EHR clinical profiles alongside time-varying dynamic environmental arrays within an attention-based multi-task framework.
Model Explainability	Captum / SHAP Core Library	Computes mathematically rigorous local feature-attributions to maintain complete clinical safety transparency for providers.
Database Tier	TimescaleDB + HIPAA-Compliant PostgreSQL Instance	Optimized storage array designed to handle highly demanding time-series climate parameters alongside relational clinical metadata.
Hosting & Security	Google Kubernetes Engine (GKE) + KMS End-to-End Encryption	Enforces full HIPAA isolation compliance; data is encrypted at rest via custom key-management keys and in transit via TLS 1.3.

9. Buildathon Dallas 2026 Strategic Blueprint (June 18–19, 2026)

Buildathon Dallas 2026 is an intense, execution-driven hackathon focused on shipping working code and robust architectures rather than basic slides. To maximize your engineering evaluation score and burn cloud credits provided by venue infrastructure sponsors, your build is divided into a strict 90/10 pipeline execution plan:

9.1 The 90% Pre-Build Framework (Construct Now)

Prior to entering the event space, complete the frontend presentation and data normalization logic locally:

1. **Frontend Application Shell (Using Lovable):** Use Lovable to write your entire user interface layout now. Create the responsive mockup of the 72-Hour Provider Triage Dashboard (complete with case management tables, risk metric sliders, and sample SHAP graph visualizations) and the matching consumer-facing Health-Weather application view, giving you a fully polished layout ready for judges.
2. **Spatial Ingestion Infrastructure:** Build your local Python mapping scripts using the h3 library. Ensure your code effortlessly handles coordinates, maps them to Resolution 7 cells, and references a local static database matching hexagons to Texas ZIP codes.
3. **Model Logic Skeletons:** Code your PyTorch TFT neural network layout and configure your fallback CatBoost pipeline layers using open-source, historical climate and demographic arrays to confirm your model input/output dimension paths are stable.
4. **Legal Entity Protection (Using SeedLegals):** Use SeedLegals to finalize your founder agreements, mutual non-disclosure configurations, and intellectual property assignments ahead of investor and recruiter evaluation.

9.2 The 10% Live-Build Framework (Execute at the Event)

Save your active cloud infrastructure orchestration and live data streaming tasks for the 32-hour live event to effectively utilize venue sponsor credits:

1. **Live Cloud & MLOps Ingestion:** Spin up your live cloud servers, microservices, and model endpoints using the infrastructure and GPU compute clusters awarded on-site by the hackathon sponsors.
2. **Real-Time Event Stream Simulation:** Configure a live, active Kafka instance at the venue. Build an active loop script that streams real-time Dallas weather and pollution API data and bends it with an event-driven synthetic stream of simulated FHIR patient clinical entries.
3. **Agentic Generation Architecture (Using MindStudio):** Connect your live backend prediction outputs to a **MindStudio AI Worker Agent**. When the multi-task model flags a patient risk delta, pass that contextual string to MindStudio to automatically generate, translate, and format an empathetic, highly tailored preventative text reminder message, ready for clinical single-click approval.